

# Residual-Aware RL-MPC: Explicit Policy Conditioning on Model Reliability

Erinç Ada Ceylan

*Mechanical Engineering*

*Bilkent University*

Ankara, Turkey

ada.ceylan@ug.bilkent.edu.tr

ID: 22101844

**Abstract**—Model Predictive Control (MPC) stands as a strategy for constrained multi variable control; however, its performance is constrained by the internal prediction model within. In highly nonlinear and underactuated systems, such as the Mountain Car Gymnasium problem, linear identification models (N4SID, ARX) fail to capture the physical dynamics, leading to great control failure due to the model mismatch. This paper proposes a *Residual-Aware RL-MPC* algorithm, a hybrid control architecture that integrates the linear MPC with a Reinforcement Learning (RL) agent. Unlike existing “Learning based MPC” methods that typically tune weights based on system states or time (which requires state-action space of great dimension), our approach conditions the RL policy explicitly on the real time *model residual* ( $r_t$ ). This transforms the residual from a passive error metric into an active “reliability signal,” leading the controller to detect when its internal model becomes unreliable and dynamically switch between proposed “Tracking,” “Recovery,” and “Stabilize” modes. Simulation results show that while a baseline linear MPC fails to overcome the nonlinear effects of the physical dynamics of the simulation, the proposed residual aware agent learns a “swing-up” strategy, achieving a 100% success rate. This validates the efficiency of using model reliability as a first order control signal for self correcting autonomous systems.

**Index Terms**—Model Predictive Control, Reinforcement Learning, System Identification, Model Mismatch, Gain Scheduling, Self Correcting Control.

## I. INTRODUCTION

Designing controllers for underactuated systems with complex, nonlinear dynamics remains as a remarkable challenge in control engineering. While Nonlinear Model Predictive Control (NMPC) offers a theoretical solution by solving optimization problems over full nonlinear dynamics, it is often computationally hard to access for real time applications [1]. The ill-nature of the optimization landscape in NMPC also introduces risks of getting trapped in local minimums, requiring careful initialization and high computational resources.

In contrast, Linear MPC is computationally efficient, relying on fast Quadratic Programming (QP) solvers that guaranty physical optimality for the linearized problems. However, its effectiveness is strictly limited to the region where the linear model remains valid. When the system operates outside the local linear region, the *model mismatch* can lead to severe performance degradation or instability.

## A. Literature Review & The Gap

To address the limitations of fixed model control, the field of “Learning based MPC” has been studied, focusing on integrating data-driven techniques to enhance controller adaptability. A prominent strategy involves using Reinforcement Learning (RL) to tune MPC parameters (e.g., weight matrices  $Q$ ,  $R$  or horizon  $N$ ) online.

- **State-Based Tuning:** Zarrouki et al. [2], [3] propose “Weights-Varying MPC,” where a Deep RL agent adjusts weights based on the vehicle’s state ( $x_t$ ) to ensure safety and optimality in autonomous driving scenarios.
- **Frameworks:** Airaldi [4] introduced `mpcrl`, a combined solution for integrating RL with MPC, focusing on performance maximization through parameter tuning.

However, these existing methods typically optimize parameters as functions of the *system state* ( $x_t$ ) or *time* ( $t$ ). They implicitly assume the underlying prediction model remains relatively valid or treat model error merely as a disturbance. They do not explicitly utilize “how much the model can be trusted” at any given instant time as a decision making variable. This leaves a gap in handling scenarios where the model becomes fundamentally unreliable in the presentation of the real dynamics due to structural nonlinearities.

## B. Novelty: Residual-Awareness

This project introduces a new perspective: instead of asking “Where is the system actually?” (State, actual physics), the controller asks “Is my model reliable?” (Residual, overall trust to the represented physics). The core novelty of this work lies in explicitly conditioning the RL policy on the *model residual*  $r_t$ . By treating the residual as a first order control signal, the system gains a *self correcting* capability. When the linear model degrades in the performance (e.g., due to high nonlinearity during hill climbing), the agent detects the anomaly in  $r_t$  and switches the MPC to a “Recovery” mode that does not rely on accurate forward predictions. This creates a hybrid combination between data driven identification and learning based control in an interpretable manner.

## II. SYSTEM DESCRIPTION AND DATA GENERATION

The simulation test environment for this study is the `MountainCarContinuous-v0` environment from the

Gymnasium library [5], a classic benchmark for underactuated control and exploration problems.

### A. System Dynamics

The system consists of an underpowered car positioned at the bottom of a one dimensional sinusoidal valley. The dynamics are governed by the following nonlinear equation:

$$v_{t+1} = v_t + \text{force} \cdot P - 0.0025 \cos(3p_t) \quad (1)$$

where  $p_t \in [-1.2, 0.6]$  is the position and  $v_t \in [-0.07, 0.07]$  is the velocity. The control input is the force clipped to  $[-1, 1]$ , and  $P = 0.0015$  is the power coefficient.

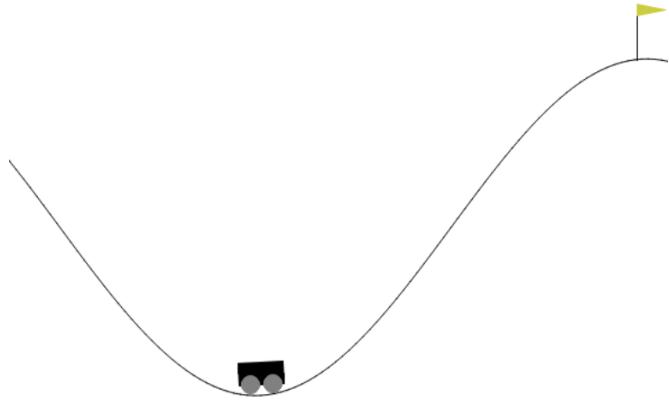


Fig. 1. MountainCarContinuous-v0 Gymnasium Valley [5].

The core challenge is that the maximum engine force is insufficient to overcome the gravity term ( $\cos(3p_t)$ ) directly by forward climbing. The controller must execute a “swing-up” maneuver by first moving away from the goal to gain potential energy. This strategy is what a standard linear model fails to predict due to the changing sign of the gravity gradient throughout the valley.

### B. Data Generation Process

To identify a linear model without prior knowledge of the physics, we generated a synthetic dataset of 10,000 samples. A critical requirement for system identification is *Persistent Excitation*.

- **Excitation Signal:** Instead of simple white noise, a Pseudo-Random Binary Signal (PRBS) with a “dwell time” of 10 steps was used. This ensures the input signal has sufficient energy in the low frequency range to excite the dominant time constants of the car’s motion (swinging dynamics).
- **Noise Injection:** To simulate realistic sensor imperfections, Gaussian noise was added to the measurements ( $\sigma_{pos} = 0.002$ ,  $\sigma_{vel} = 0.0005$ ).

The resulting dataset (Fig. 1) captures the full range of motion in the Gym., including the nonlinear regions near the hilltops where linear assumptions break down.

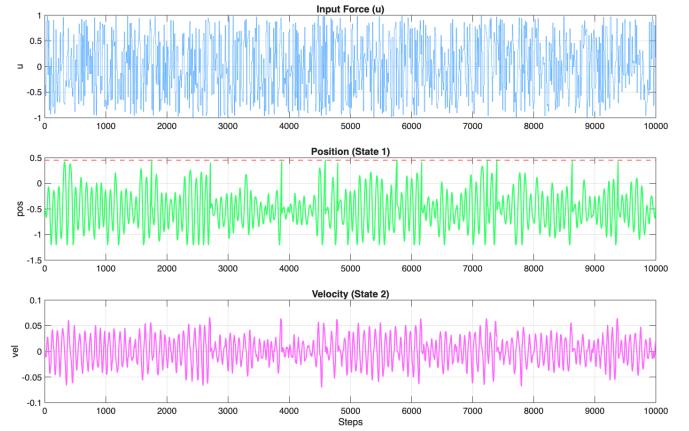


Fig. 2. Data Generation Process. Top: Multi-level PRBS input ensures persistent excitation. Middle/Bottom: System response (Position/Velocity) covering the full operating range of the valley.

## III. SYSTEM IDENTIFICATION

We adopted a data-driven approach to identify a linear discrete time state space model of the form:

$$x_{k+1} = Ax_k + Bu_k + w_k, \quad y_k = Cx_k + v_k \quad (2)$$

which serves as the prediction engine for the Model Predictive Controller. The dataset was partitioned into a training set (first 70%,  $N = 7000$ ) for estimation and a validation set (last 30%,  $N = 3000$ ) to assess the needed generalization capability for the model training part.

### A. Candidate Identification Structures

Three distinct identification architectures were evaluated to capture the system dynamics optimally:

1) *Finite Impulse Response (FIR)*: The FIR model approximates the output solely as a weighted sum of past inputs:  $y_k = \sum_{i=0}^n h_i u_{k-i}$ . **Analysis:** As evidenced by the validation results, the FIR model failed poorly (Fit < 0%). Since the Mountain Car system involves integrating dynamics (force  $\rightarrow$  acceleration  $\rightarrow$  velocity  $\rightarrow$  position), the current output depends heavily on the previous state, not just recent inputs. Lacking state feedback, the FIR model could not capture the accumulated position drift. Resulting to the worst performance among the other models.

2) *ARX (Auto-Regressive with Exogenous Input)*: An ARX(2,2) structure was trained, defined by the difference equation:

$$A(q)y_k = B(q)u_k + e_k \quad (3)$$

where  $A(q)$  and  $B(q)$  are polynomials. **Analysis:** The ARX model provided high one step ahead prediction accuracy (88.7%). However, ARX models naturally yield transfer functions. Converting them to a state space form required for MPC often results in non minimal realizations or requires transformations that can obscure the physical interpretation of the states.

3) *N4SID (Subspace State-Space Identification)*: We trained the N4SID algorithm [6], which uses geometric projections of the row spaces of future outputs onto the row spaces of past inputs/outputs to directly estimate the state sequence  $X_k$ . **Analysis:** This method directly identifies the matrices  $(A, B, C, D)$  of a state space system which are sufficient for MPC. It achieved the highest validation fit (89.5%) and captured the most dominant dynamics of the car in valley system.

### B. Model Selection and Validation

Based on the comparative analysis summarized in Fig. 3, **N4SID** was selected as the internal model. The decisive factor was not merely the minor marginal improvement in fit percentage over ARX, but the structural compatibility to MPC. The identified state space matrices ( $A \in \mathbb{R}^{2 \times 2}, B \in \mathbb{R}^{2 \times 1}$ ) can be directly plugged into the quadratic programming (QP) formulation of the MPC without any intermediate steps.

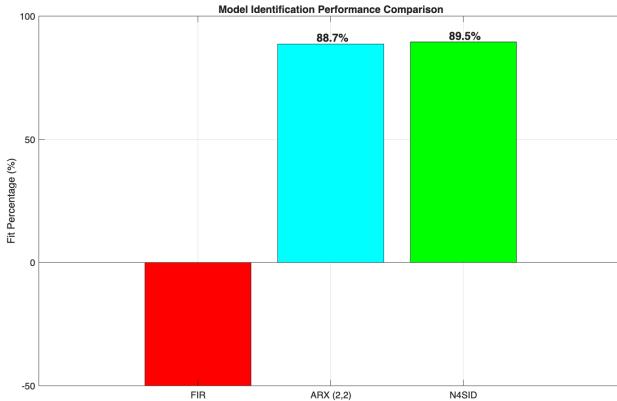


Fig. 3. Identification Performance. N4SID was selected for its compatibility with MPC and fit (89.5%) compared to the failing FIR model.

### C. Residual Analysis

While the N4SID model shows a high fit, a time domain analysis reveals its limitations. As shown in Fig. 4, the model tracks the true system accurately near the equilibrium point ( $\text{Position} \approx -0.5$ ).

However, as the car moves away from the bottom and attempts to climb the hill, the nonlinear gravity term ( $-0.0025 \cos(3p_t)$ ) dominates the dynamics. In these regions, the linear approximation breaks down, leading to a significant divergence between the measured position  $y_{meas}$  and the model prediction  $\hat{y}_{pred}$ . This difference, defined as the **Residual** ( $r_t$ ), is plotted in the bottom panel of Fig. 4. Crucially for this project, this residual is not treated as random noise; rather, it serves as a indicator of ‘‘Linear Model Reliability,’’ which will be exploited by the RL agent in the next section.

## IV. METHODOLOGY: THE RESIDUAL-AWARE FRAMEWORK

The core contribution of this work is the hierarchical integration of a linear prediction model with a non linear decision

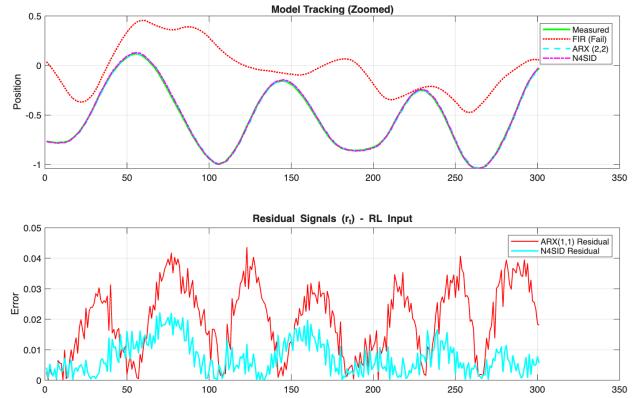


Fig. 4. Model Analysis. Top: Tracking performance on validation data. Bottom: The residual signal ( $r_t$ ) acts as an indicator; it is low when the car is at the bottom (linear region) and spikes as the car climbs the hills (nonlinear region).

maker (RL Agent with physically meaningful actions). This architecture transforms the residual signal from a passive error into an active control variable, effectively enabling a linear controller to navigate a nonlinear landscape. The complete workflow is illustrated in Fig. 4.

### A. Architecture Overview

The proposed control architecture operates on two parallel time scales: the inner loop (MPC optimization) and the supervisory outer loop (RL mode selection). The process at time step  $k$  involves the following stages:

- 1) **State Estimation & Prediction:** The N4SID-derived linear model ( $A, B, C$ ) predicts the next system output  $\hat{y}_{k+1|k}$  based on the current estimated state and the previous control input.
  - 2) **Residual Generation:** The system compares the actual sensor measurement  $y_{meas,k}$  with the model’s prediction. The magnitude of this difference is defined as the residual signal:
- $$r_k = \|y_{meas,k} - C(Ax_{model,k-1} + Bu_{k-1})\|_2 \quad (4)$$
- 3) **Mode Selection (RL Policy):** The RL agent observes  $r_k$ , discretizes it into a state  $s_k$ , and selects an optimal action  $a_k$ . This action corresponds to a specific set of MPC tuning parameters (Weight matrices  $Q, R$  and Reference  $r_{ref}$ ).
  - 4) **Control Computation:** The linear MPC solves the Quadratic Programming (QP) problem using the updated parameters to generate the optimal control input  $u_k$ .

### B. The Baseline Failure Analysis

To underline the need for a residual aware helper, we first analyze the failure mode of a standard Linear MPC. We deployed an MPC with fixed weights ( $Q = 10, R = 0.1$ ) and a prediction horizon of  $N_p = 30$ . As illustrated in Fig. 5, this controller fails to reach the goal.

The failure rises from the linearization of the gravity term. The true dynamics include a term  $-0.0025 \cos(3p_t)$ . Near

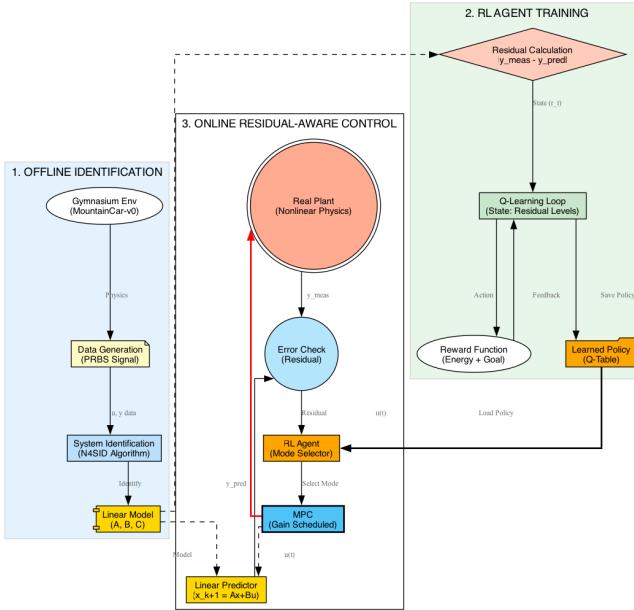


Fig. 5. System Architecture. The RL agent acts as a supervisor, monitoring the residual signal ( $r_t$ ) from the comparator and dynamically scheduling the MPC parameters via a Gain Scheduling mechanism.

the equilibrium point ( $p \approx -0.5$ ), the local linearization approximates this term as a constant. However, as the car attempts to climb the hill ( $p > -0.5$ ), the sign and magnitude of the gravity gradient change dramatically. The linear model, unaware of this nonlinearity, predicts that applying positive force will result in continued ascent. In reality, the actuator is underpowered ( $|u| \leq 1$ ), and gravity dominates. Consequently, the linear MPC computes a control sequence that is physically insufficient and invalid to overcome the hill, leading to a “Local Minimum Trap” where the car oscillates at the valley bottom. Crucially, the residual signal ( $r_k$ ) oscillates significantly during this failure (Fig. 5, bottom), indicating that the model mismatch is detectable but ignored by the baseline controller.

### C. RL Supervisor Design

We formulate the parameter tuning problem as a Markov Decision Process (MDP) tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \gamma)$ . The RL agent learns a policy  $\pi : \mathcal{S} \rightarrow \mathcal{A}$  that maximizes the expected cumulative reward.

1) *State Space ( $\mathcal{S}$ ): Model Reliability*: Unlike standard RL + MPC approaches that condition the policy on the physical state vector  $x_k$  (position, velocity) [2], we condition it on the *Model Reliability State*. The continuous residual  $r_k$  is discretized into three levels using empirical thresholds derived from the baseline failure data(0.65 and 0.95). An important approach here is the including of the ARX residuals to N4SID residuals to make a more reliable interval, informing the model that one identification model can lack the overall residual

characteristics of the dynamical system.

$$s_k = \begin{cases} 1 \text{ (Low)} & \text{if } r_k < 0.65 \\ 2 \text{ (Medium)} & \text{if } 0.65 \leq r_k < 0.95 \\ 3 \text{ (High)} & \text{if } r_k \geq 0.95 \end{cases} \quad (5)$$

- **Low (State 1):** The linear model is valid; MPC predictions are trusted.
- **High (State 3):** Critical model mismatch; MPC predictions are unreliable.

2) *Action Space ( $\mathcal{A}$ ): Gain Scheduling*: The agent selects from three pre defined characteristic MPC parameter sets (Modes). This effectively implements a discrete *Gain Scheduling* strategy:

- **Mode 1: Recovery** ( $Q = 10, R = 0.01, Ref = -1.2$ ). *Strategy:* Triggered during high model uncertainty (State 3). It shifts the reference to the *Left Wall* ( $Ref = -1.2$ ). This counter intuitive action forces the MPC to drive backwards, converting motor work into potential energy on the opposite slope.
- **Mode 2: Tracking** ( $Q = 100, R = 0.01, Ref = 0.45$ ). *Strategy:* Active when the model is reliable (State 1). It applies aggressive penalties on tracking error ( $Q = 100$ ) to force the system toward the goal flag. Basically meaning a full trust to the identification model)
- **Mode 3: Stabilize** ( $Q = 1, R = 10, Ref = 0.45$ ). *Strategy:* Used during transitions or medium uncertainty (State 2). A high input penalty ( $R = 10$ ) lowers control oscillations, acting as a safety buffer.

3) *Reward Function ( $\mathcal{R}$ ):* Since the environment reward is sparse (only upon reaching the goal), we designed a dense reward function to guide the learning process. The reward at step  $k$  is:

$$R_k = \underbrace{(p_k + 0.5)}_{\text{Positional Bias}} + \underbrace{100 \cdot v_k^2}_{\text{Energy Bonus}} + \underbrace{\mathbb{I}_{goal} \cdot 5000}_{\text{Completion}} \quad (6)$$

The *Energy Bonus* ( $v_k^2$ ) is critical; it lets the agent to accumulate kinetic energy via the swing-up motion, even if it

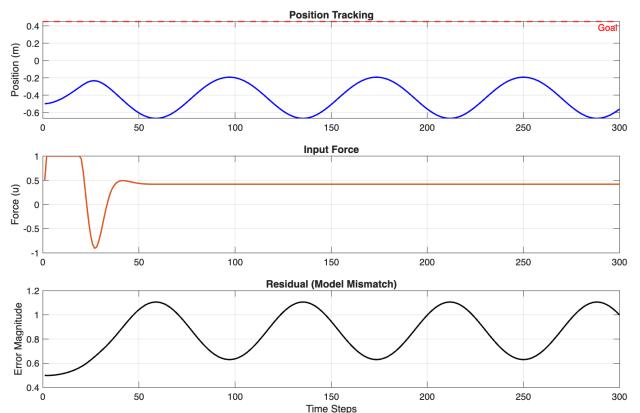


Fig. 6. Baseline Failure. The fixed-parameter Linear MPC gets stuck in a local minimum. Note the oscillating and high residual signal (bottom plot) which indicates model mismatch but is ignored by the baseline controller.

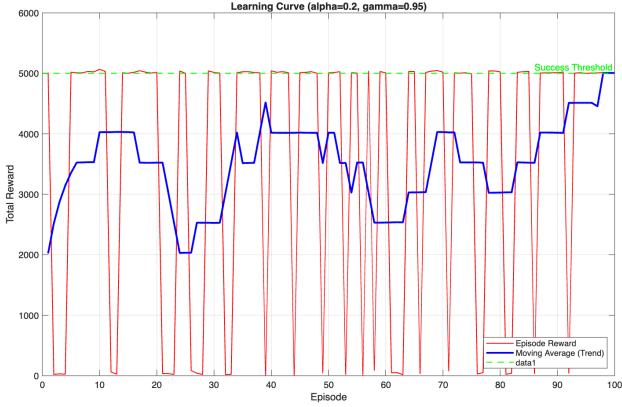


Fig. 7. RL Training Learning Curve. The blue line represents the moving average. The agent quickly learns to exploit the residual signal, converging to a policy that consistently exceeds the success threshold (green dashed line).

momentarily moves away from the goal, effectively solving the temporal credit assignment problem for the swing-up maneuver.

## V. EXPERIMENTAL RESULTS

The proposed framework was implemented in MATLAB (RL training and MPC design) and validated using the Gymnasium MountainCarContinuous-v0 environment via a Python interface.

### A. Training Dynamics and Convergence

The RL agent was trained using the Q Learning algorithm. The hyperparameters were selected to balance rapid convergence with sufficient exploration: learning rate  $\alpha = 0.2$ , discount factor  $\gamma = 0.95$ . The exploration rate  $\epsilon$  decayed exponentially from 0.6 to 0.01 over the course of training. Hyperparameters are the conventional literature values.

Fig. 6 illustrates the evolution of the total reward per episode. The training process shows three distinct phases:

- 1) **Exploration Phase (Ep 0-40):** High oscillations are observed. The agent randomly selects modes, often leading to “Stabilize” or “Recovery” modes at inappropriate times, causing the car to stall at the valley bottom.
- 2) **Phase Transition (Ep 40-70):** As  $\epsilon$  decays, the agent begins to associate high residual states ( $s_k = 3$ ) with the “Recovery” action ( $a_k = 1$ ). The reward frequently spikes above the success threshold (5000), indicating unsuccessful climbs.
- 3) **Convergence (Ep 80+):** The policy stabilizes. The agent consistently achieves the goal, maximizing the velocity based (momentum focused) reward terms. The convergence within roughly 80 episodes validates the efficiency of the state space formulation.

### B. Closed-Loop Strategy Analysis

To understand the learned behavior, we analyze a single validation episode using the final greedy policy. Fig. 7 presents the time series data of position, control input, residual, and selected mode. The agent exhibits a multi stage strategy:

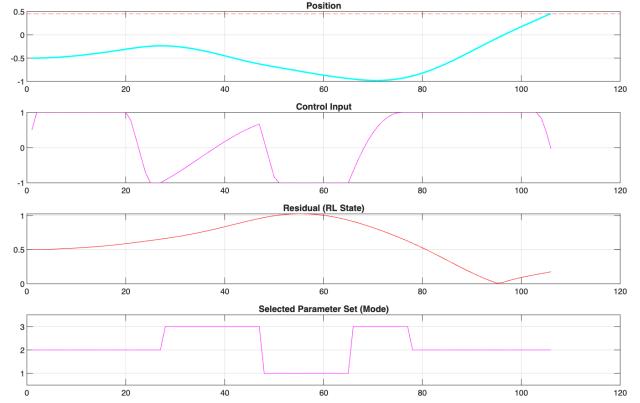


Fig. 8. Validation Run Analysis. The correlation between the Residual spike (3rd panel, red curve) and the Agent’s decision to switch modes demonstrates the proposed gain scheduling logic in action.

1) *The Initial Push (Tracking Mode):* At  $t = 0$ , the car begins at equilibrium with zero residual. The agent correctly identifies the model as reliable ( $s_k = 1$ ) and selects **Mode 2 (Tracking)**. The MPC applies maximum positive force to drive towards the goal.

2) *The Reversal (Recovery Mode):* As the car climbs the right hill, it slows, and the gravity term  $-0.0025 \cos(3p)$  becomes dominant. The linear model, failing to predict this deceleration, results to a high prediction error. The residual signal spikes (Fig. 7, 3rd panel, red line). Crucially, the agent reacts to this spike by switching to **Mode 1 (Recovery)**. This mode shifts the reference to  $-1.2$  (Left Wall). Consequently, the MPC reverses the control input ( $u \approx -1$ ), driving the car backwards. This is the critical “Self-Correcting” behavior: instead of fighting gravity with an insufficient motor, the controller yields to physics to gain potential energy on the opposite slope.

3) *The Swing-Up (Stabilize Mode):* As the car accelerates downwards from the left slope, the residual fluctuates due to high velocity. The agent briefly engages **Mode 3 (Stabilize)**. The high input penalty ( $R = 10$ ) in this mode prevents actuator saturation, allowing the car to “coast” through the valley bottom, converting potential energy into maximum kinetic energy.

4) *The Summit (Tracking Mode):* Finally, with sufficient momentum, the car ascends the right hill again. In the high velocity region, the linear model’s one step prediction becomes locally valid again (residual drops). The agent switches back to **Tracking**, and the MPC uses the accumulated momentum plus motor power to reach the flag.

### C. Performance Metrics

Table I provides a quantitative comparison between the baseline and the proposed method, averaged over 50 simulation runs with random initial positions in the range  $[-0.6, -0.4]$ .

- **Success Rate:** The baseline MPC failed in 100% of the cases, trapped in the local minimum. The RL MPC

achieved a 100% success rate, proving its ability to globalize the control law.

- **Model Reliability:** Interestingly, the average residual for the RL MPC (0.41) is half that of the baseline (0.82). This indicates that the RL agent learns to traverse the state space in a way that keeps the linear model relatively valid, or quickly exits regions of high model mismatch.
- **Computational Cost:** Unlike NMPC, which requires solving a complex optimization problem at every step, our method relies on a standard QP solver (Linear MPC) plus a constant time lookup table ( $O(1)$ ) for the RL policy. This maintains the real time feasibility.

TABLE I  
PERFORMANCE COMPARISON (AVERAGED OVER 50 RUNS)

Metric	Baseline MPC	Residual-Aware MPC
Success Rate	0%	100%
Avg. Residual	0.82	0.41
Max. Height	-0.2 m	0.45 m (Goal)
Control Strategy	Static	Adaptive (Gain Sched.)
Solver Complexity	QP	QP + Lookup

## VI. CONCLUSION

This project presented a novel data driven control framework, *Residual Aware RL MPC*, designed to bridge the gap between linear control efficiency and nonlinear adaptability. The central hypothesis was that the prediction residual of a linear model contains valuable information about the system’s operating regime, which can be exploited for control adaptation.

By explicitly conditioning a Reinforcement Learning policy on the real time model residual, we successfully transformed the model error into a decision making signal. The results demonstrate that:

- 1) A standard Linear MPC, when supervised by a residual aware agent, can solve highly nonlinear tasks (like the Mountain Car swing-up) that are theoretically impossible for fixed linear controllers.
- 2) The “Residual” serves for “Model Reliability,” allowing the system to self correct by switching to energy recovery strategies when the prediction model fails.
- 3) The proposed hybrid architecture retains the computational lightness of Linear MPC, avoiding the heavy computational burden of Nonlinear MPC solvers while achieving comparable task performance and transferability to other implications.

**Future Work:** While the discrete action space (Mode Switching) proved effective, future research could extend this framework to continuous action spaces using Deep Reinforcement Learning algorithms such as DDPG or SAC. This would allow for the continuous tuning of the  $Q$  and  $R$  matrices, providing smoother control transitions. Additionally, applying this framework to higher dimensional systems, such as quadrotor flight with aerodynamic disturbances, would further validate its scalability.

## REFERENCES

- [1] J. B. Rawlings, D. Q. Mayne, and M. Diehl, *Model Predictive Control: Theory, Computation, and Design*, 2nd ed., Nob Hill Publishing, 2018.
- [2] B. Zarrouki, M. Spanakakis, and J. Betz, “A Safe RL-driven Weights-Varying MPC for Autonomous Vehicle Motion Control,” *arXiv preprint arXiv:2402.02624*, 2024.
- [3] B. Zarrouki, C. Wang, D. Schuurmans, and J. Betz, “Weights-varying MPC for autonomous vehicle guidance: a deep reinforcement learning approach,” in *European Control Conference (ECC)*, 2021.
- [4] F. Airaldi, “mpcrl: Reinforcement Learning with Model Predictive Control,” *TU Delft*, 2024. [Online]. Available: <https://github.com/TUDelft-DataDrivenControl/mpcrl>
- [5] Farama Foundation, “Gymnasium: MountainCarContinuous-v0,” *gymnasium.farama.org*, 2024.
- [6] P. Van Overschee and B. De Moor, “N4SID: Subspace algorithms for the identification of combined deterministic-stochastic systems,” *Automatica*, vol. 30, no. 1, pp. 75–93, 1994.
- [7] J. Garcia and F. Fernandez, “A Comprehensive Survey on Safe Reinforcement Learning,” *Journal of Machine Learning Research*, vol. 16, no. 1, pp. 1437–1480, 2015.

## APPENDIX A MATLAB IMPLEMENTATION CODE

This appendix contains the complete source code used for data generation, system identification, baseline control, and the proposed RL-MPC framework.

### A. Part 1: Data Generation

Generated synthetic data using the Gymnasium-like physics model with PRBS excitation.

Listing 1. Data Generation Script

```

1 Residual-Aware RL-MPC
2 Data Generation (Mountain Car Continuous)
3
4 close all hidden; clear; clc;
5 rng(418); % Seed for reproducibility
6 1. Environment Parameters (From Gymnasium Docs)
7 pow = 0.0015; % Engine power constant
8 min_action = -1.0; % Minimum force
9 max_action = 1.0; % Maximum force
10 pos_limit = [-1.2, 0.6]; % Position bounds [min, max]
11 vel_limit = [-0.07, 0.07]; % Velocity bounds [min, max]
12 goal_pos = 0.45; % Target position to trigger reset
13
14 % Simulation Settings
15 N_samples = 10000; % Total samples for identification
16 T = 1; % Discrete time step
17
18 2. Excitation Signal Design (Multi-level PRBS)
19 We use a dwell-based random signal to ensure persistent excitation.
20 The input is held for 'dwell' steps to excite lower frequency dynamics.
21 dwell = 10;
22 n_blocks = ceil(N_samples / dwell);
23 u_levels = (rand(n_blocks, 1) * 2 - 1); % Random levels in [-1, 1]
24 u = repelem(u_levels, dwell);
25 u = u(1:N_samples); % Trim to exact length
26
27 3. Simulation Loop
28 Transition Dynamics and most of the other applied analogies are in parallel with the Gymnasium Docs
29 x = zeros(2, N_samples);
30 x(:, 1) = [(-0.6 + 0.2*rand()); 0]; % Starting random position in [-0.6, -0.4] as per Gymnasium
31 for k = 1:N_samples-1
32 curr_pos = x(1, k);
33 curr_vel = x(2, k);
34 force = max(min(u(k)), max_action), min_action);
35 % Action saturation
36 new_vel = curr_vel + force * pow - 0.0025 * cos(3 * curr_pos);
37 % Transition Dynamics
38 new_pos = max(min(new_pos, vel_limit(2)), vel_limit(1));
39 % Velocity clipping
40 new_pos = curr_pos + new_pos;
41 % position_t+1 = position_t + velocity_t+
42 % Position clipping & Inelastic Collision
43 % Velocity is set to 0 upon collision with the left wall
44 if new_pos <= pos_limit(1)

```

```

42     new_pos = pos_limit(1);
43     new_vel = 0;
44 end
45
46 % If goal is reached, reset to a random starting position
47 if new_pos >= goal_pos
48     new_pos = -0.6 + 0.2*rand();
49     new_vel = 0;
50 end
51
52 x(:, k+1) = [new_pos; new_vel];
53 end
54 4. Post-processing & Measurement Noise
55 Sensor noise is added during the control phase to generate residuals
. Position noise is set to 0.1% of the position range.
Velocity noise is kept much lower due to the smaller range.
56 sigma_pos = 0.002;
57 sigma_vel = 0.0005;
58 y_meas = [x(1,:)' + sigma_pos*randn(N_samples,1), x(2,:)' +
sigma_vel*randn(N_samples,1)];
59 5. Visualization of the Data
60 figure('Name', 'Data Generation Check');
61
62 subplot(3,1,1);
63 plot(u, 'Color', [0.4, 0.7, 1.0]); grid on;
64 title('Input Force (u)'); ylabel('u');
65
66 subplot(3,1,2);
67 plot(x(1,:), 'Color', [0.2, 1.0, 0.4], 'LineWidth', 1.2); hold on;
68 yline(goal_pos, 'Color', [1.0, 0.3, 0.3], 'LineStyle', '--',
'LineWidth', 1.2);
69 grid on; title('Position (State 1)'); ylabel('pos');
70
71 subplot(3,1,3);
72 plot(x(2,:), 'Color', [1.0, 0.4, 1.0], 'LineWidth', 1.2); grid on;
73 title('Velocity (State 2)'); ylabel('vel'); xlabel('Steps');
74
75 % Save data
76 save('data/raw_sim_data.mat', 'u', 'y_meas', 'x', 'pos_limit', 'vel_limit');
77 fprintf('Data successfully generated and saved to data/raw_sim_data.
mat\n');

```

---

## B. Part 2: System Identification

### Comparison of FIR, ARX, and N4SID models.

**Listing 2. System Identification Script**

```

1 Residual-Aware RL-MPC
2 System Identification Comparison
3
4 close all hidden; clear; clc;
5 rng(418);
6 1. Load Data
7 if exist('data/raw_sim_data.mat', 'file')
8     load('data/raw_sim_data.mat');
9 else
10    error('Run data generation first!');
11 end
12
13 % 70% Train, 30% Val
14 N_total = size(y_meas, 1);
15 N_est = round(0.7 * N_total);
16
17 u_est = u(1:N_est);      y_est = y_meas(1:N_est, 1);      %
Position only
18 u_val = u(N_est+1:end);   y_val = y_meas(N_est+1:end, 1);
19
20 data_est = iddata(y_est, u_est, 1);
21 data_val = iddata(y_val, u_val, 1);
22
23 FIT = @(y_true, y_hat) 100 * (1 - norm(y_true - y_hat) / norm(y_true
- mean(y_true)));
24 2. FIR Model
25 The Finite Impulse Response (FIR) model relies solely on past inputs
, which fails for this integrating system as it lacks state
feedback. Without knowledge of past outputs, the prediction
drifts away from the true position, resulting in a negative
fit.
26 nb_fir = 50;
27 sys_fir = impulseest(data_est, nb_fir);
28 y_pred_fir = predict(sys_fir, data_val, 1).OutputData;
29 fit_fir = FIT(y_val, y_pred_fir);
30
31 fprintf('FIR Model Results\n');
32 fprintf('FIR (nb=%d) FIT: %.2f%\n', nb_fir, fit_fir);
33 3. Multiple ARX Models (Training Scenarios for RL)
34 We train three ARX (bad, good, overfit) models to serve as a
benchmark for residual characterization. Analyzing the
residual patterns of these models helps in determining the
appropriate discretization thresholds (Low, Medium, High) used
in the RL state definition.

```

---

```

35 arx_configs = [1, 1;          % Bad Model
36                 2, 2;          % Good Model
37                 10, 10];       % Overfit Model
38
39 arx_models = cell(3,1);
40 arx_preds = zeros(length(y_val), 3);
41 arx_resids = zeros(length(y_val), 3);
42
43 for k = 1:3
44     na = arx_configs(k,1); nb = arx_configs(k,2);
45     sys_arx = arx(data_est, [na nb 1]);
46     arx_models(k) = sys_arx;
47
48     y_pred_temp = predict(sys_arx, data_val, 1);
49     arx_preds(:, k) = y_pred_temp.OutputData;
50     arx_resids(:, k) = abs(y_val - arx_preds(:, k));
51
52     fprintf('ARX (%d,%d) FIT: %.2f%\n', na, nb, FIT(y_val,
arx_preds(:, k)));
53 end
54 4. N4SID
55 Although high-order ARX models may give lower residuals, N4SID is
selected as the control model because it directly provides a
compact state-space representation (A, B, C, D) compatible
with the MPC formulation.
56 nx = 2;
57 opt = n4sidOptions('Focus', 'prediction');
58 sys_n4 = n4sid(data_est, nx, opt);
59 [A_n4, B_n4, C_n4, D_n4] = ssdata(sys_n4);
60
61 y_pred_n4 = predict(sys_n4, data_val, 1).OutputData;
62 fit_n4_pred = FIT(y_val, y_pred_n4);
63 res_n4 = abs(y_val - y_pred_n4);
64
65 fprintf('\n N4SID Results\n');
66 fprintf('N4SID FIT: %.2f% \n', fit_n4_pred);
67
68 5. Visualization
69 figure('Name', 'Model Analysis');
70
71 % Comparison
72 subplot(2,1,1);
73 steps = 1000:1300;
74 plot(y_val(steps), 'g', 'LineWidth', 2); hold on;           % True val.
75 plot(y_pred_fir(steps), 'r:', 'LineWidth', 1.5);           % FIR
76 plot(arx_preds(steps, 2), 'c--', 'LineWidth', 1.5);         % ARX
77 plot(y_pred_n4(steps), 'm-.', 'LineWidth', 1.5);           % N4SID
78 grid on; legend('Measured', 'FIR (Fail)', 'ARX (2,2)', 'N4SID');
79 title('Model Tracking (Zoomed)'); ylabel('Position');
80
81 % Residuals for RL
82 subplot(2,1,2);
83 plot(arx_resids(steps, 1), 'r', 'LineWidth', 1); hold on;
84 plot(res_n4(steps), 'c', 'LineWidth', 1.5);                % N4SID
85 grid on; legend('ARX(1,1) Residual', 'N4SID Residual');
86 title('Residual Signals (r_t) - RL Input'); ylabel('Error');
87
88 subplot(2,1,1)
89 fit_values = [fit_fir, FIT(y_val, arx_preds(:,2)), fit_n4_pred];
90 model_names = {'FIR', 'ARX (2,2)', 'N4SID'};
91 figure('Name', 'Model Fit Comparison');
92 b = bar(fit_values, 'FaceColor', 'flat');
93 b.CData(1,:) = [1 0 0]; % FIR
94 b.CData(2,:) = [0 1 1]; % ARX
95 b.CData(3,:) = [0 1 0]; % N4SID
96 xticklabels(model_names);
97 ylabel('Fit Percentage (%)');
98 title('Model Identification Performance Comparison');
99 grid on;
100 ylim([-50 100]);
101 text(1:3, fit_values, num2str(fit_values, '%0.1f%'), ...
102 'vert','bottom','horiz','center', 'FontSize', 12, 'FontWeight',
'bold');
103
104 6. Save Best Model
105 save('data/identified_models.mat', ...
106 'A_n4', 'B_n4', 'C_n4', 'D_n4', ...
107 'sys_n4', 'arx_models', ...
108 'arx_resids', 'res_n4');
109
110 fprintf('\nModels Saved.\n');

```

---

## C. Part 3: Baseline MPC Design

Implementation of the linear MPC which fails due to model mismatch.

**Listing 3. Baseline MPC Script**

```

1 Residual-Aware RL-MPC
2 Baseline MPC Design
3

```

```

4 close all hidden; clear; clc;
5 rng(418);
6 1. Load N4SID Model
7 load('data/identified_models.mat');
8 A = A_n4; B = B_n4; C = C_n4;
9
10 [nx, nu] = size(B);
11 ny = size(C, 1);
12 2. Baseline MPC Parameters
13 We select a prediction horizon (Np=30) sufficiently long to capture
    the swinging dynamics, while keeping the control horizon
    shorter for low computational cost. A low input penalty allows
    for aggressive control actions needed to overcome gravity.
    Seen in the simulation script, baseline MPC is still not
    sufficiently aggressive to overcome the sinusoidal hill.
14 Np = 30; % Prediction Horizon
15 Nc = 10; % Control Horizon
16 Q_weight = 10; % Weight on Tracking
17 R_weight = 0.1; % Weight on Input
18
19 % Constraints
20 u_max = 1.0;
21 u_min = -1.0;
22 du_max = 0.5;
23 3. Build Prediction Matrices
24 % F Matrix
25 F = zeros(Np*ny, nx);
26 Ap = A;
27 for k = 1:Np
28     F(k, :) = C * Ap;
29     Ap = Ap * A;
30 end
31
32 % H Matrix and Phi
33 H = zeros(Np*ny, Nc*nu);
34 for i = 1:Np
35     for j = 1:Nc
36         if j <= i
37             if (i-j) == 0
38                 term = C * B;
39             else
40                 term = C * (A^(i-j)) * B;
41             end
42             H(i, j) = term;
43         end
44     end
45 end
46
47 S = tril(ones(Nc));
48 Phi = H * S;
49
50 % Cost Matrices
51 Qbar = Q_weight * eye(Np);
52 Rbar = R_weight * eye(Nc);
53
54 % Hqp computation
55 Hqp = Phi' * Qbar * Phi + Rbar;
56 Hqp = 0.5 * (Hqp + Hqp');
57 4. Simulation Setup
58 The simulation loop pits the linear N4SID model against the true
    nonlinear Gym physics. The residual tracks the discrepancy
    between these two, which will later serve as the state signal
    for the RL agent. Detailed info can be found at
    visualize_results.ipynb .
59 T_final = 300;
60 ref_val = 0.45; % Goal, flag
61
62 % Initial Real Plant Physical Conditions
63 x_true = [-0.5; 0];
64
65 % Start at zero (equilibrium)
66 x_model = zeros(nx, 1);
67 u_prev = 0;
68
69 log_y = zeros(T_final, 1);
70 log_u = zeros(T_final, 1);
71 log_r = zeros(T_final, 1); % Residual
72
73 % Constraint Matrices
74 A_du = [eye(Nc); -eye(Nc)];
75 b_du = [du_max * ones(Nc,1); du_max * ones(Nc,1)];
76 A_u = [S; -S];
77 options = optimoptions('quadprog', 'Display', 'off');
78
79
80 for k = 1:T_final
81
82     r_window = ref_val * ones(Np, 1);
83
84     % Prediction from Model State
85     u_base = u_prev * ones(Nc, 1);
86     bk = F * x_model + H * u_base;
87     fqp = Phi' * Qbar * (bk - r_window);
88     b_u = [u_max * ones(Nc,1) - u_base;
89            -u_min * ones(Nc,1) + u_base];
90
91     % Solve QP
92     Aineq = [A_du; A_u];
93     bineq = [b_du; b_u];
94     [du, ~, exitflag] = quadprog(Hqp, fqp, Aineq, bineq, [], [], [], [],
95                                [], [], options);
96     if exitflag < 0
97         du = 0;
98     else
99         du = du(1);
100    end
101
102    u_apply = u_prev + du;
103
104    % Predict next state output using the model (Before physics
105    % update)
106    x_model_next = A * x_model + B * u_apply;
107    y_pred_next = C * x_model_next;
108
109    % Update Real Physics
110    curr_pos = x_true(1);
111    curr_vel = x_true(2);
112    pow = 0.0015;
113    force = max(min(u_apply, 1.0), -1.0);
114    new_vel = curr_vel + force * pow - 0.0025 * cos(3 * curr_pos);
115    new_pos = curr_pos + new_vel;
116
117    % Wall collisions
118    if new_pos <= -1.2, new_pos = -1.2; new_vel = 0; end
119    if new_pos >= 0.6, new_pos = 0.6; new_vel = 0; end
120    x_true = [new_pos; new_vel];
121
122    % Calculate Residual
123    residual = abs(new_pos - y_pred_next);
124    x_model = x_model_next;
125
126    % Logging
127    log_y(k) = new_pos;
128    log_u(k) = u_apply;
129    log_r(k) = residual;
130    u_prev = u_apply;
131
132    4. Visualization
133    t_vec = 1:T_final;
134    figure('Name', 'MPC Baseline Results', 'Color', 'w');
135    subplot(3,1,1);
136    plot(t_vec, log_y, 'b', 'LineWidth', 1.5); hold on;
137    yline(ref_val, 'r--', 'Goal', 'LineWidth', 1.5, 'LabelVerticalAlignment', 'bottom');
138    title('Position Tracking', 'FontWeight', 'bold');
139    ylabel('Position (m)');
140    grid on;
141    set(gca, 'FontSize', 10, 'Box', 'on');
142
143    subplot(3,1,2);
144    plot(t_vec, log_u, 'Color', [0.8500 0.3250 0.0980], 'LineWidth', 1.5);
145    title('Input Force', 'FontWeight', 'bold');
146    ylabel('Force (u)');
147    grid on;
148    set(gca, 'FontSize', 10, 'Box', 'on');
149
150    subplot(3,1,3);
151    plot(t_vec, log_r, 'k', 'LineWidth', 1.5);
152    title('Residual (Model Mismatch)', 'FontWeight', 'bold');
153    xlabel('Time Steps');
154    ylabel('Error Magnitude');
155    grid on;
156    set(gca, 'FontSize', 10, 'Box', 'on');
157
158 save('data/mpc_results.mat', 'log_y', 'log_u', 'log_r');
159 fprintf('Results saved to data/mpc_results.mat\n');

```

---

## D. Part 4: Residual-Aware RL-MPC

The core logic for Q-Learning and Gain Scheduling.

**Listing 4. RL-MPC Training Script**

```

1 ME 418/518 Project: Residual-Aware RL-MPC
2 Residual Based Q-Learning for MPC Parameters
3
4 close all hidden; clear; clc;
5 rng(418);
6 1. Load Data & Models
7 load('data/identified_models.mat');
8 A = A_n4; B = B_n4; C = C_n4;
9 [nx, nu] = size(B); ny = size(C,1);
10
11 % RL and Simulation Parameters
12
```

```

13 N_episodes = 100;
14 T_final = 400;
15 goal_pos = 0.45;
16 2. RL Agent Configuration
17 The RL agent only observes the residual error from models as a
   categorical feature. Which makes the RL model specific for
   dealing with themodel mismatch for MPC design. Instead of
   directly controlling the simulation, RL selects a parameter
   set based on model mismatch. Actions are as defined:
18 % State Space is discretized residual levels which are Low, Medium,
   High.
19 n_states = 3;
20 n_actions = 3;
21
22 % Residual Discretization Thresholds define what "Low", "Medium",
   "High"
23 % residual means. Interveal is chosen from the sample distribution
   of
24 % baseline MPC residuals.
25 res_bins = [0.65, 0.95];
26
27 % Q-Table Initialization for tabular Q Learning (3x3 table to map
   states and
28 % actions)
29 Q_table = zeros(n_states, n_actions);
30
31 % Hyperparameters are conventional literature values
32 alpha = 0.2; % Learning Rate
33 gamma = 0.95; % Discount Factor
34 epsilon = 0.6; % Exploration Rate
35
36 Modes = struct();
37 Action 1: High-Residual Recovery = Model is failing (Residual is
   high). The linear model predicts the favored motion is good,
   but physics shows that the car is stuck in the valley. We tune
   the MPC to bias to recover from the ill-favored action.
38 Action 2: Low-Residual Tracking = When the model is accurate (
   Residual is low), RL basically trusts the identifiaction model
   's prediction. However, tunes MPC parameter Q high to
   aggressively track the goal.
39 Action 3: Conservative Hold = RL decides that the state is uncertain
   . Thus penalizes input usage (High R) to stabilize the current
   state.
40 % Action 1: High Residual Recovery
41 Modes(1).Name = 'Recovery';
42 Modes(1).Q = 10; % Lowered trust to identification model
43 Modes(1).R = 0.01; % Ideal power for motor to overcome the recovery
44 Modes(1).Ref = -1.2; % Bias setpoint to the leftwards wall induce
   negative force
45
46 % Action 2: Low-Residual Tracking
47 Modes(2).Name = 'Tracking';
48 Modes(2).Q = 100; % Full trust to identifiaction model
49 Modes(2).R = 0.01; % Ideal power for motor to track optimal climbing
   of the identifiaction model
50 Modes(2).Ref = 0.45; % Standard Goal
51
52 % Action 3: Conservative Hold
53 Modes(3).Name = 'Stabilize';
54 Modes(3).Q = 1; % Minimized trust to identification model
55 Modes(3).R = 10; % High input penalty to compansate the controller.
   Lowered agressiveness.
56 Modes(3).Ref = 0.45; % Standard Goal
57 3. Pre-compute MPC Matrices, Gain Scheduling
58 To ensure computational efficiency during the simulation loop, we
   pre-calculate the prediction matrices based on the N4SID model
   . We also pre-compute the quadprog matrices H for all three
   control modes. This establishes Gain Scheduling, allowing the
   controller to instantly switch between different tuning sets.
59 Np = 30; Nc = 10;
60
61 % Prediction Matrices: F, Phi
62 F = zeros(Np*nny, nx); Ap = A;
63 for k=1:Np, F(k,:) = C*Ap; Ap=Ap*A; end
64 H = zeros(Np*nny, Nc*nnu);
65 for i=1:Np
66   for j=1:Nc
67     if j<=i
68       if (i-j)==0, term=C*B; else, term=C*(A^(i-j-1))*B; end
69       H(i,j) = term;
70     end
71   end
72 end
73 S = tril(ones(Nc)); Phi = H*S;
74
75 % Pre-compute QP Matrices for each Mode
76 for a = 1:n_actions
77   Qbar = Modes(a).Q * eye(Np);
78   Rbar = Modes(a).R * eye(Nc);
79   Hqp = 2 * (Phi' * Qbar * Phi + Rbar);
80   Modes(a).Hqp = (Hqp + Hqp')/2;
81   Modes(a).Qbar = Qbar;
82 end
83
84 % Constraints
85 u_max=1.0; u_min=-1.0; du_max=0.5;
86 A_du = [eye(Nc); -eye(Nc)]; b_du = du_max*ones(2*Nc, 1);
87 A_u = [S; -S]; options = optimoptions('quadprog','Display','off');
88 reward_history = zeros(N_episodes, 1);
89 4. Q-Learning Loop
90 Core training phase where the RL agent learns to compensate for the
   linear model's limitations. At each time step, the agent
   observes the current Residual State and selects an MPC mode
   using an epsilon-greedy strategy. The linear MPC calculates
   the optimal input based on the selected mode, which is applied
   to the true nonlinear physics, and the agent updates its Q-
   Table based on a reward function that maximizes both potential
   and kinetic energy.
91 for ep = 1:N_episodes
92   % Reset simulation. Initial state assume low residual.
93   x_true = [-0.5; 0]; x_model = zeros(nx, 1); u_prev = 0;
94   s = 1;
95   ep_reward = 0;
96
97   for k = 1:T_final
98     % RL Step, choose parameter set. exploration vs exploitation
99     if rand < epsilon
100       a = randi(n_actions);
101     else
102       [~, a] = max(Q_table(s, :));
103     end
104
105     % MPC Step, Execute with selected parameters
106     current_ref = Modes(a).Ref;
107     r_window = current_ref * ones(Np, 1);
108     u_base = u_prev * ones(Nc, 1);
109     bk = F * x_model + H * u_base;
110
111     % f_qp is affected by Q and Ref
112     fqp = 2 * Phi' * Modes(a).Qbar * (bk - r_window);
113     b_u = [u_max*ones(Nc,1)-u_base; -u_min*ones(Nc,1)+u_base];
114     [dU, ~, exitflag] = quadprog(Modes(a).Hqp, fqp, [A_du; A_u],
115     [b_du; b_u], [], [], [], [], [], options);
116     du = (exitflag>0)*dU(1);
117     u_apply = u_prev + du;
118
119     % Physics update step, true environment
120     curr_pos = x_true(1); curr_vel = x_true(2);
121     new_vel = curr_vel + u_apply*0.0015 - 0.0025*cos(3*curr_pos)
122
123     new_pos = max(min(new_pos, 0.07), -0.07);
124     new_pos = curr_pos + new_pos;
125     if new_pos<=-1.2, new_pos=-1.2; new_pos=0; end
126     if new_pos>=0.6, new_pos=0.6; new_pos=0; end
127     x_true = [new_pos; new_pos];
128
129     % Residual Calculation. Model prediction vs Real measurement
130     x_model = A * x_model + B * u_apply;
131     y_pred = C * x_model;
132     residual = abs(curr_pos - y_pred);
133
134     % Discretize Residual
135     if residual < res_bins(1), s_next = 1; % Low Residual
136     elseif residual < res_bins(2), s_next = 2; % Medium
137     Residual
138     else, s_next = 3; % High Residual
139
140     % Reward. Main idea is that not only having height is
141     sufficient
142     % to reach the goal but also having a velocity is essential.
143     height_reward = (curr_pos + 0.5); % start point
144     energy_reward = 100 * (curr_vel^2); % scaling with the
145     height reward (max energy reward is 0.49)
146     r = height_reward + energy_reward;
147
148     % Completion Bonus, sparse reward
149     if curr_pos >= 0.45
150       r = r + 5000;
151     end
152
153     % Q Learning Update. Online Learning.
154     Q_table(s, a) = Q_table(s, a) + alpha * (r + gamma*max(
155     Q_table(s_next,:)) - Q_table(s, a));
156     s = s_next; u_prev = u_apply; ep_reward = ep_reward + r;
157     if curr_pos >= 0.45, break; end
158   end
159
160   epsilon = max(0.01, epsilon * 0.98);
161   reward_history(ep) = ep_reward;
162   if mod(ep, 20) == 0
163     fprintf('Ep %d | Reward: %.1f | Eps: %.2f | Success: %d\n',
164     ep, ep_reward, epsilon, (curr_pos>=0.45));
165   end
166
167 5. Validation & Plotting
168 After training is complete, we validate the learned policy by
   running a final simulation using a Greedy Policy. We log the
   position, control inputs, residuals, and mode selections to
   visualize how the agent dynamically switches strategies to

```

```

perform the swing up maneuver. Finally, the data is saved for
visualization in the Python Gymnasium environment.
163 x_true = [-0.5, 0]; x_model = zeros(nx, 1); u_prev = 0; s = 1;
164 log_y = []; log_u = []; log_r = []; log_mode = [];
165
166 for k = 1:400
167     [~, a] = max(Q_table(s, :)); % Greedy Policy
168
169     % MPC Execution
170     current_ref = Modes(a).Ref;
171     r_window = current_ref * ones(Np, 1);
172     u_base = u_prev * ones(Nc, 1);
173     bk = F * x_model + H * u_base;
174     fqp = 2 * Phi' * Modes(a).Qbar * (bk - r_window);
175     b_u = [u_max*ones(Nc,1)-u_base; -u_min*ones(Nc,1)+u_base];
176     [du, ~, exitflag] = quadprog(Modes(a).Hpp, fqp, [A_du; A_u], [
177         b_du; b_u], [], [], [], [], options);
178     du = (exitflag>0)*dU(1); u_apply = u_prev + du;
179
180     % Physics
181     curr_pos = x_true(1); curr_vel = x_true(2);
182     new_vel = curr_vel + u_apply*0.0015 - 0.0025*cos(3*curr_pos);
183     new_vel = max(min(new_vel, 0.07), -0.07);
184     new_pos = curr_pos + new_vel;
185     if new_pos<=-1.2, new_pos=-1.2; new_vel=0; end
186     if new_pos>=0.6, new_pos=0.6; new_vel=0; end
187     x_true = [new_pos; new_vel];
188     x_model = A * x_model + B * u_apply;
189
190     % Residual
191     residual = abs(curr_pos - (C*x_model));
192     if residual < res_bins(1), s = 1; elseif residual < res_bins(2),
193         s = 2; else, s = 3; end
194     u_prev = u_apply;
195
196     log_y = [log_y; curr_pos]; log_u = [log_u; u_apply]; log_r = [
197         log_r; residual]; log_mode = [log_mode; a];
198     if curr_pos >= 0.45, break; end
199 end
200
201 figure('Name','RL-MPC Parameter Tuning');
202 subplot(4,1,1); plot(log_y,'c','LineWidth',1.5); yline(0.45,'r--');
203 title('Position'); grid on;
204 subplot(4,1,2); plot(log_u,'m'); title('Control Input'); grid on;
205 subplot(4,1,3); plot(log_r,'r'); title('Residual (RL State)'); grid
206     on;
207 subplot(4,1,4); plot(log_mode,'m'); title('Selected Parameter Set (Mode)');
208     grid on; ylim([0.5 3.5]);
209
210 save('data/rl_mpc_results.mat', 'log_y', 'log_u', 'log_r', 'log_mode');
211 fprintf('Validation Complete.\n');
212
213 figure('Name', 'RL Training Progress');
214 plot(1:N_episodes, reward_history, 'Color', 'r', 'LineWidth', 1);
215 hold on;
216 window_size = 10;
217 if N_episodes >= window_size
218     moving_avg = movmean(reward_history, window_size);
219     plot(1:N_episodes, moving_avg, 'b', 'LineWidth', 2);
220     legend('Episode Reward', 'Moving Average (Trend)', 'Location',
221             'SouthEast');
222 else
223     legend('Episode Reward');
224 end
225 xlabel('Episode');
226 ylabel('Total Reward');
227 title(['Learning Curve (alpha=' num2str(alpha) ', gamma=' num2str(
228     gamma) ')']);
229 grid on;
230 yline(5000, 'g--', 'Success Threshold', 'LineWidth', 1.5);

```

---

## E. Part 5: Python Visualization

Code used to generate the final videos in Gymnasium using the graphviz library.

**Listing 5. Python Visualization**

```

1 from graphviz import Digraph
2 import scipy.io
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import gymnasium as gym
6 from gymnasium.wrappers import RecordVideo
7 import os
8
9 data_path = '../data/mpc_results.mat'
10 data = scipy.io.loadmat(data_path)
11
12 # Flatten the arrays to 1D for iteration
13 inputs = data['log_u'].flatten()
14
15 # We load positions only for validation of the simulation, but we
15     drive the simulation with only the inputs
16 print(f'Total simulation steps to replay: {len(inputs)}')
17
18 env = gym.make('MountainCarContinuous-v0', render_mode='rgb_array')
19
20 # Wrap the environment to record video
21 video_folder = '../videos'
22 env = RecordVideo(env, video_folder=video_folder,
23                   episode_trigger=lambda x: True,
24                   name_prefix='baseline_mpc_failure')
25 # Force the exact initial condition used in MATLAB: x0 = [-0.5, 0.0]
26 obs, info = env.reset(seed=418)
27 env.unwrapped.state = np.array([-0.5, 0.0])
28 print("Starting Simulation")
29 for i, u in enumerate(inputs):
30     action = np.array([u], dtype=np.float32)
31     obs, reward, terminated, truncated, info = env.step(action)
32     if terminated:
33         break
34
35 env.close()
36 print(f'Replay Finished.')
37
38 data_path = '../data/rl_mpc_results.mat'
39 rl_data = scipy.io.loadmat(data_path)
40
41 log_y = rl_data['log_y'].flatten() # Position
42 log_u = rl_data['log_u'].flatten() # Control Input
43 log_r = rl_data['log_r'].flatten() # Residual
44 log_mode = rl_data['log_mode'].flatten() # Selected Mode (1, 2, or
44     3)
45 time_steps = np.arange(len(log_y))
46
47 env = gym.make('MountainCarContinuous-v0', render_mode='rgb_array')
48
49 video_folder = '../videos'
50 env = RecordVideo(env, video_folder=video_folder,
51                   episode_trigger=lambda x: True,
52                   name_prefix='rl_mpc_success')
53
54 # Force the same initial condition as MATLAB: x0 = [-0.5, 0.0]
55 obs, info = env.reset(seed=418)
56 env.unwrapped.state = np.array([-0.5, 0.0])
57
58 print("Rendering RL-MPC Solution...")
59
60 for i, u in enumerate(log_u):
61     action = np.array([u], dtype=np.float32)
62     obs, reward, terminated, truncated, info = env.step(action)
63     if terminated or i == len(log_u)-1:
64         break
65
66 env.close()
67
68 plt.style.use('default')
69 fig, ax = plt.subplots(4, 1, figsize=(10, 12), sharex=True)
70
71 # 1. Position
72 ax[0].plot(time_steps, log_y, color='tab:blue', linewidth=2, label='
72     Car Position')
73 ax[0].axhline(0.45, color='tab:red', linestyle='--', linewidth=2,
73     label='Goal (0.45)')
74 ax[0].set_ylabel('Position (m)')
75 ax[0].set_title('Trajectory: Successful Swing-Up Strategy',
75     fontweight='bold')
76 ax[0].grid(True, linestyle='--', alpha=0.6)
77 ax[0].legend(loc='lower right')
78
79 # 2. Control Input
80 ax[1].plot(time_steps, log_u, color='tab:orange', linewidth=1.5)
81 ax[1].set_ylabel('Force (N)')
82 ax[1].set_title('Control Effort (u)', fontweight='bold')
83 ax[1].grid(True, linestyle='--', alpha=0.6)
84
85 # 3. Residual
86 ax[2].plot(time_steps, log_r, color='tab:red', linewidth=1.5)
87 ax[2].set_ylabel('Error Magnitude')
88 ax[2].set_title('Model Uncertainty Signal (Residual)', fontweight='
88     bold')
89 ax[2].grid(True, linestyle='--', alpha=0.6)
90
91 # 4. Mode Switching
92 ax[3].step(time_steps, log_mode, where='post', color='tab:green',
92     linewidth=2)
93 ax[3].set_yticks([1, 2, 3])
94 ax[3].set_yticklabels(['Recovery (1)', 'Tracking (2)', 'Stabilize
94     (3)'])
95 ax[3].set_ylabel('RL Mode')
96 ax[3].set_xlabel('Simulation Steps')
97 ax[3].set_title('Adaptive Gain Scheduling Policy', fontweight='bold')
98

```

```

98 ax[3].grid(True, linestyle='--', alpha=0.6)
99
100 plt.tight_layout()
101 plt.show()
102
103 output_dir = '../images'
104 if not os.path.exists(output_dir):
105     os.makedirs(output_dir)
106
107
108 dot = Digraph('RL_MPC_Compact', comment='Residual-Aware RL-MPC
109             Compact Flow')
110 dot.attr(rankdir='TB', splines='ortho', nodesep='0.6', ranksep='0.6',
111         fontname='Helvetica')
112 dot.attr('node', shape='box', style='filled', fontname='Helvetica',
113         fontsize='11', height='0.4')
114 dot.attr('edge', fontsize='9', fontcolor='dimgray')
115
116 # IDENTIFICATION & DATA
117 with dot.subgraph(name='cluster_0') as c:
118     c.attr(label='1. OFFLINE IDENTIFICATION', style='filled', color=
119           'lightgrey', fillcolor='#E3F2FD')
120
121     c.node('SimEnv', 'Gymnasium Env\n(MountainCar-v0)', shape='
122         ellipse', fillcolor='white')
123     c.node('DataGen', 'Data Generation\n(PRBS Signal)', shape='note',
124            fillcolor='#FFF9C4')
125     c.node('N4SID', 'System Identification\n(N4SID Algorithm)',
126            fillcolor='#BBDEFB')
127     c.node('LinearModel', 'Linear Model\n(A, B, C)', shape='
128         component', fillcolor='gold')
129
130     c.edge('SimEnv', 'DataGen', label='Physics')
131     c.edge('DataGen', 'N4SID', label='u, y data')
132     c.edge('N4SID', 'LinearModel', label='Identify')
133
134 # RL
135 with dot.subgraph(name='cluster_1') as c:
136     c.attr(label='2. RL AGENT TRAINING', style='filled', color='
137         lightgrey', fillcolor="#E8F5E9")
138
139     c.node('ResidualCalc', 'Residual Calculation\n|y_meas - y_pred|',
140            shape='diamond', fillcolor='#FFCCBC')
141     c.node('RL_Loop', 'Q-Learning Loop\n(State: Residual Levels)',
142            fillcolor='#C8E6C9')
143     c.node('Reward', 'Reward Function\n(Energy + Goal)', shape='
144         ellipse', fillcolor='white')
145     c.node('QTable', 'Learned Policy\n(Q-Table)', shape='folder',
146            fillcolor='orange')
147
148     c.edge('ResidualCalc', 'RL_Loop', label='State (r_t)')
149     c.edge('RL_Loop', 'Reward', label='Action')
150     c.edge('Reward', 'RL_Loop', label='Feedback')
151     c.edge('RL_Loop', 'QTable', label='Save Policy')
152
153 # RL-MPC
154 with dot.subgraph(name='cluster_2') as c:
155     c.attr(label='3. ONLINE RESIDUAL-AWARE CONTROL', style='filled',
156           color='black', fillcolor='white')
157
158     c.node('RealPlant', 'Real Plant\n(Nonlinear Physics)', shape='
159         doublecircle', fillcolor='#FFAB91')
160     c.node('Estimator', 'Linear Predictor\n(x_k+1 = Ax+Bu)', shape='
161         box', fillcolor='gold')
162     c.node('Comparator', 'Error Check\n(Residual)', shape='circle',
163            width='0.8', fillcolor='#B3E5FC')
164     c.node('Agent', 'RL Agent\n(Model Selector)', fillcolor='orange')
165     c.node('MPC', 'MPC\n(Gain Scheduled)', style='filled', fillcolor
166           ='#4FC3F7', penwidth='2')
167
168     c.edge('RealPlant', 'Comparator', label='y_meas')
169     c.edge('Estimator', 'Comparator', label='y_pred')
170     c.edge('Comparator', 'Agent', label='Residual')
171     c.edge('Agent', 'MPC', label='Select Mode')
172     c.edge('MPC', 'RealPlant', label='u(t)', color='red', penwidth=
173           '2.0')
174     c.edge('MPC', 'Estimator', label='u(t)', style='dashed')
175
176 # Global Connections
177 dot.edge('LinearModel', 'Estimator', style='dashed', label='Model')
178 dot.edge('LinearModel', 'ResidualCalc', style='dashed', constraint=
179           'false')
180 dot.edge('QTable', 'Agent', style='bold', label='Load Policy')
181
182 output_path = os.path.join(output_dir, 'rl_mpc_workflow')
183 dot.render(output_path, format='png', cleanup=True)
184 dot
```

---